

Building External Modules

This document describes how to build an out-of-tree kernel module.

=== Table of Contents

- === 1 Introduction
- === 2 How to Build External Modules
 - 2.1 Command Syntax
 - 2.2 Options
 - 2.3 Targets
 - 2.4 Building Separate Files
- === 3. Creating a Kbuild File for an External Module
 - 3.1 Shared Makefile
 - 3.2 Separate Kbuild file and Makefile
 - 3.3 Binary Blobs
 - 3.4 Building Multiple Modules
- === 4. Include Files
 - 4.1 Kernel Includes
 - 4.2 Single Subdirectory
 - 4.3 Several Subdirectories
- === 5. Module Installation
 - 5.1 INSTALL_MOD_PATH
 - 5.2 INSTALL_MOD_DIR
- === 6. Module Versioning
 - 6.1 Symbols From the Kernel (vmlinux + modules)
 - 6.2 Symbols and External Modules
 - 6.3 Symbols From Another External Module
- === 7. Tips & Tricks
 - 7.1 Testing for CONFIG_FOO_BAR

=== 1. Introduction

"kbuild" is the build system used by the Linux kernel. Modules must use kbuild to stay compatible with changes in the build infrastructure and to pick up the right flags to "gcc." Functionality for building modules both in-tree and out-of-tree is provided. The method for building either is similar, and all modules are initially developed and built out-of-tree.

Covered in this document is information aimed at developers interested in building out-of-tree (or "external") modules. The author of an external module should supply a makefile that hides most of the complexity, so one only has to type "make" to build the module. This is easily accomplished, and a complete example will be presented in section 3.

=== 2. How to Build External Modules

To build external modules, you must have a prebuilt kernel available that contains the configuration and header files used in the build. Also, the kernel must have been built with modules enabled. If you are using a distribution kernel, there will be a package for the kernel you are running provided by your distribution.

An alternative is to use the "make" target "modules_prepare." This will make sure the kernel contains the information required. The target exists solely as a simple way to prepare a kernel source tree for building external modules.

NOTE: "modules_prepare" will not build Module.symvers even if CONFIG_MODULEVERSIONS is set; therefore, a full kernel build needs to be executed to make module versioning work.

--- 2.1 Command Syntax

The command to build an external module is:

```
$ make -C <path_to_kernel_src> M=$PWD
```

The kbuild system knows that an external module is being built due to the "M=<dir>" option given in the command.

To build against the running kernel use:

```
$ make -C /lib/modules/`uname -r`/build M=$PWD
```

Then to install the module(s) just built, add the target "modules_install" to the command:

```
$ make -C /lib/modules/`uname -r`/build M=$PWD modules_install
```

--- 2.2 Options

(\$KDIR refers to the path of the kernel source directory.)

```
make -C $KDIR M=$PWD
```

-C \$KDIR

The directory where the kernel source is located. "make" will actually change to the specified directory when executing and will change back when finished.

M=\$PWD

Informs kbuild that an external module is being built. The value given to "M" is the absolute path of the directory where the external module (kbuild file) is located.

--- 2.3 Targets

When building an external module, only a subset of the "make" targets are available.

```
make -C $KDIR M=$PWD [target]
```

The default will build the module(s) located in the current directory, so a target does not need to be specified. All output files will also be generated in this directory. No attempts are made to update the kernel source, and it is a precondition that a successful "make" has been executed for the kernel.

modules

The default target for external modules. It has the same functionality as if no target was specified. See description above.

modules_install

Install the external module(s). The default location is /lib/modules/<kernel_release>/extra/, but a prefix may be added with INSTALL_MOD_PATH (discussed in section 5).

clean

Remove all generated files in the module directory only.

help

List the available targets for external modules.

--- 2.4 Building Separate Files

It is possible to build single files that are part of a module. This works equally well for the kernel, a module, and even for external modules.

Example (The module `foo.ko`, consist of `bar.o` and `baz.o`):

```
make -C $KDIR M=$PWD bar.lst
make -C $KDIR M=$PWD baz.o
make -C $KDIR M=$PWD foo.ko
make -C $KDIR M=$PWD /
```

=== 3. Creating a Kbuild File for an External Module

In the last section we saw the command to build a module for the running kernel. The module is not actually built, however, because a build file is required. Contained in this file will be the name of the module(s) being built, along with the list of requisite source files. The file may be as simple as a single line:

```
obj-m := <module_name>.o
```

The kbuild system will build `<module_name>.o` from `<module_name>.c`, and, after linking, will result in the kernel module `<module_name>.ko`. The above line can be put in either a "Kbuild" file or a "Makefile." When the module is built from multiple sources, an additional line is needed listing the files:

```
<module_name>-y := <src1>.o <src2>.o ...
```

NOTE: Further documentation describing the syntax used by kbuild is located in `Documentation/kbuild/makefiles.txt`.

The examples below demonstrate how to create a build file for the module `8123.ko`, which is built from the following files:

```
8123_if.c
8123_if.h
8123_pci.c
8123_bin.o_shipped      <= Binary blob
```

--- 3.1 Shared Makefile

An external module always includes a wrapper makefile that supports building the module using "make" with no arguments. This target is not used by kbuild; it is only for convenience. Additional functionality, such as test targets, can be included but should be filtered out from kbuild due to possible name clashes.

Example 1:

```
--> filename: Makefile
ifneq ($(KERNELRELEASE),)
# kbuild part of makefile
obj-m := 8123.o
8123-y := 8123_if.o 8123_pci.o 8123_bin.o

else
# normal makefile
KDIR ?= /lib/modules/`uname -r`/build

default:
$(MAKE) -C $(KDIR) M=$$PWD

# Module specific targets
genbin:
echo "X" > 8123_bin.o_shipped

endif
```

The check for `KERNELRELEASE` is used to separate the two parts of the makefile. In the example, kbuild will only see the two assignments, whereas "make" will see everything except these two assignments. This is due to two passes made on the file:

the first pass is by the "make" instance run on the command line; the second pass is by the kbuild system, which is initiated by the parameterized "make" in the default target.

--- 3.2 Separate Kbuild File and Makefile

In newer versions of the kernel, kbuild will first look for a file named "Kbuild," and only if that is not found, will it then look for a makefile. Utilizing a "Kbuild" file allows us to split up the makefile from example 1 into two files:

Example 2:

```
--> filename: Kbuild
obj-m := 8123.o
8123-y := 8123_if.o 8123_pci.o 8123_bin.o

--> filename: Makefile
KDIR ?= /lib/modules/`uname -r`/build

default:
    $(MAKE) -C $(KDIR) M=$$PWD

# Module specific targets
genbin:
    echo "X" > 8123_bin.o_shipped
```

The split in example 2 is questionable due to the simplicity of each file; however, some external modules use makefiles consisting of several hundred lines, and here it really pays off to separate the kbuild part from the rest.

The next example shows a backward compatible version.

Example 3:

```
--> filename: Kbuild
obj-m := 8123.o
8123-y := 8123_if.o 8123_pci.o 8123_bin.o

--> filename: Makefile
ifneq ($(KERNELRELEASE),)
# kbuild part of makefile
include Kbuild

else
# normal makefile
KDIR ?= /lib/modules/`uname -r`/build

default:
    $(MAKE) -C $(KDIR) M=$$PWD

# Module specific targets
genbin:
    echo "X" > 8123_bin.o_shipped

endif
```

Here the "Kbuild" file is included from the makefile. This allows an older version of kbuild, which only knows of makefiles, to be used when the "make" and kbuild parts are split into separate files.

--- 3.3 Binary Blobs

Some external modules need to include an object file as a blob. kbuild has support for this, but requires the blob file to be named <filename>_shipped. When the kbuild rules kick in, a copy of <filename>_shipped is created with _shipped stripped off, giving us <filename>. This shortened filename can be used in the assignment to the module.

Throughout this section, `8123_bin.o_shipped` has been used to build the kernel module `8123.ko`; it has been included as `8123_bin.o`.

```
8123-y := 8123_if.o 8123_pci.o 8123_bin.o
```

Although there is no distinction between the ordinary source files and the binary file, `kbuild` will pick up different rules when creating the object file for the module.

--- 3.4 Building Multiple Modules

`kbuild` supports building multiple modules with a single build file. For example, if you wanted to build two modules, `foo.ko` and `bar.ko`, the `kbuild` lines would be:

```
obj-m := foo.o bar.o
foo-y := <foo_srcs>
bar-y := <bar_srcs>
```

It is that simple!

=== 4. Include Files

Within the kernel, header files are kept in standard locations according to the following rule:

- * If the header file only describes the internal interface of a module, then the file is placed in the same directory as the source files.
- * If the header file describes an interface used by other parts of the kernel that are located in different directories, then the file is placed in `include/linux/`.

NOTE: There are two notable exceptions to this rule: larger subsystems have their own directory under `include/`, such as `include/scsi`; and architecture specific headers are located under `arch/${ARCH}/include/`.

--- 4.1 Kernel Includes

To include a header file located under `include/linux/`, simply use:

```
#include <linux/module.h>
```

`kbuild` will add options to "gcc" so the relevant directories are searched.

--- 4.2 Single Subdirectory

External modules tend to place header files in a separate `include/` directory where their source is located, although this is not the usual kernel style. To inform `kbuild` of the directory, use either `ccflags-y` or `CFLAGS_<filename>.o`.

Using the example from section 3, if we moved `8123_if.h` to a subdirectory named `include`, the resulting `kbuild` file would look like:

```
--> filename: Kbuild
obj-m := 8123.o

ccflags-y := -Iinclude
8123-y := 8123_if.o 8123_pci.o 8123_bin.o
```

Note that in the assignment there is no space between `-I` and

the path. This is a limitation of kbuild: there must be no space present.

--- 4.3 Several Subdirectories

kbuild can handle files that are spread over several directories. Consider the following example:

```
.
|__ src
|   |__ complex_main.c
|   |__ hal
|       |__ hardwareif.c
|       |__ include
|       |   |__ hardwareif.h
|__ include
|   |__ complex.h
```

To build the module complex.ko, we then need the following kbuild file:

```
--> filename: Kbuild
obj-m := complex.o
complex-y := src/complex_main.o
complex-y += src/hal/hardwareif.o

ccflags-y := -I$(src)/include
ccflags-y += -I$(src)/src/hal/include
```

As you can see, kbuild knows how to handle object files located in other directories. The trick is to specify the directory relative to the kbuild file's location. That being said, this is NOT recommended practice.

For the header files, kbuild must be explicitly told where to look. When kbuild executes, the current directory is always the root of the kernel tree (the argument to "-C") and therefore an absolute path is needed. \$(src) provides the absolute path by pointing to the directory where the currently executing kbuild file is located.

=== 5. Module Installation

Modules which are included in the kernel are installed in the directory:

```
/lib/modules/$(KERNELRELEASE)/kernel/
```

And external modules are installed in:

```
/lib/modules/$(KERNELRELEASE)/extra/
```

--- 5.1 INSTALL_MOD_PATH

Above are the default directories but as always some level of customization is possible. A prefix can be added to the installation path using the variable INSTALL_MOD_PATH:

```
$ make INSTALL_MOD_PATH=/frodo modules_install
=> Install dir: /frodo/lib/modules/$(KERNELRELEASE)/kernel/
```

INSTALL_MOD_PATH may be set as an ordinary shell variable or, as shown above, can be specified on the command line when calling "make." This has effect when installing both in-tree and out-of-tree modules.

--- 5.2 INSTALL_MOD_DIR

External modules are by default installed to a directory under `/lib/modules/$(KERNELRELEASE)/extra/`, but you may wish to locate modules for a specific functionality in a separate directory. For this purpose, use `INSTALL_MOD_DIR` to specify an alternative name to "extra."

```
$ make INSTALL_MOD_DIR=gandalf -C $KDIR \
    M=$PWD modules_install
=> Install dir: /lib/modules/$(KERNELRELEASE)/gandalf/
```

=== 6. Module Versioning

Module versioning is enabled by the `CONFIG_MODVERSIONS` tag, and is used as a simple ABI consistency check. A CRC value of the full prototype for an exported symbol is created. When a module is loaded/used, the CRC values contained in the kernel are compared with similar values in the module; if they are not equal, the kernel refuses to load the module.

`Module.symvers` contains a list of all exported symbols from a kernel build.

--- 6.1 Symbols From the Kernel (vmlinux + modules)

During a kernel build, a file named `Module.symvers` will be generated. `Module.symvers` contains all exported symbols from the kernel and compiled modules. For each symbol, the corresponding CRC value is also stored.

The syntax of the `Module.symvers` file is:

```
<CRC>          <Symbol>          <module>

0x2d036834  scsi_remove_host  drivers/scsi/scsi_mod
```

For a kernel build without `CONFIG_MODVERSIONS` enabled, the CRC would read `0x00000000`.

`Module.symvers` serves two purposes:

- 1) It lists all exported symbols from vmlinux and all modules.
- 2) It lists the CRC if `CONFIG_MODVERSIONS` is enabled.

--- 6.2 Symbols and External Modules

When building an external module, the build system needs access to the symbols from the kernel to check if all external symbols are defined. This is done in the `MODPOST` step. `modpost` obtains the symbols by reading `Module.symvers` from the kernel source tree. If a `Module.symvers` file is present in the directory where the external module is being built, this file will be read too. During the `MODPOST` step, a new `Module.symvers` file will be written containing all exported symbols that were not defined in the kernel.

--- 6.3 Symbols From Another External Module

Sometimes, an external module uses exported symbols from another external module. `kbuild` needs to have full knowledge of all symbols to avoid spitting out warnings about undefined symbols. Three solutions exist for this situation.

NOTE: The method with a top-level `kbuild` file is recommended but may be impractical in certain situations.

Use a top-level `kbuild` file

If you have two modules, `foo.ko` and `bar.ko`, where `foo.ko` needs symbols from `bar.ko`, you can use a common top-level `kbuild` file so both modules are compiled in the same build. Consider the following

directory layout:

```
./foo/ <= contains foo.ko
./bar/ <= contains bar.ko
```

The top-level kbuild file would then look like:

```
#!/Kbuild (or ./Makefile):
    obj-y := foo/ bar/
```

And executing

```
$ make -C $KDIR M=$PWD
```

will then do the expected and compile both modules with full knowledge of symbols from either module.

Use an extra Module.symvers file

When an external module is built, a Module.symvers file is generated containing all exported symbols which are not defined in the kernel. To get access to symbols from bar.ko, copy the Module.symvers file from the compilation of bar.ko to the directory where foo.ko is built. During the module build, kbuild will read the Module.symvers file in the directory of the external module, and when the build is finished, a new Module.symvers file is created containing the sum of all symbols defined and not part of the kernel.

Use "make" variable KBUILD_EXTRA_SYMBOLS

If it is impractical to copy Module.symvers from another module, you can assign a space separated list of files to KBUILD_EXTRA_SYMBOLS in your build file. These files will be loaded by modpost during the initialization of its symbol tables.

=== 7. Tips & Tricks

--- 7.1 Testing for CONFIG_FOO_BAR

Modules often need to check for certain CONFIG_ options to decide if a specific feature is included in the module. In kbuild this is done by referencing the CONFIG_ variable directly.

```
#fs/ext2/Makefile
obj-$(CONFIG_EXT2_FS) += ext2.o

ext2-y := balloc.o bitmap.o dir.o
ext2-$(CONFIG_EXT2_FS_XATTR) += xattr.o
```

External modules have traditionally used "grep" to check for specific CONFIG_ settings directly in .config. This usage is broken. As introduced before, external modules should use kbuild for building and can therefore use the same methods as in-tree modules when testing for CONFIG_ definitions.